

9/15/25, 11:01 AM

tensorflow_in_notebooks.ipynb - Colab

import tensorflow as tf

import datetime, os

TensorBoard in notebooks

Download the [FashionMNIST](#) dataset and scale it:

```
fashion_mnist = tf.keras.datasets.fashion_mnist

(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/29515/29515\_0s\_0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/26421880/26421880\_0s\_0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/5148/5148\_0s\_0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/4422102/4422102\_0s\_0us/step
```

Create a very simple model:

```
def create_model():
    return tf.keras.models.Sequential([
        tf.keras.layers.Flatten(input_shape=(28, 28), name='layers_flatten'),
        tf.keras.layers.Dense(512, activation='relu', name='layers_dense'),
        tf.keras.layers.Dropout(0.2, name='layers_dropout'),
        tf.keras.layers.Dense(10, activation='softmax', name='layers_dense_2')
    ])

Train the model using Keras and the TensorBoard callback:
```

```
def train_model():
    model = create_model()
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

    logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
    tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=

    model.fit(x=x_train,
              y=y_train,
              epochs=5,
              validation_data=(x_test, y_test),
              callbacks=[tensorboard_callback])
```

tensorflow_in_notebooks.ipynb - Colab

https://colab.research.google.com/github/tensorflow/tensorboard/blob/master/docs/tensorboard_in_notebooks.ipynb#scrollTo=TB0wBWfcVqHz&p... 2/7

9/15/25, 11:01 AM

tensorflow_in_notebooks.ipynb - Colab

Copyright 2019 The TensorFlow Authors.

Using TensorBoard in Notebooks

View on TensorFlow...

Run in Google Colab

Download notebook...

View source on GitHub...

TensorBoard can be used directly within notebook experiences such as [Colab](#) and [Jupyter](#). This can be helpful for sharing results, integrating TensorBoard into existing workflows, and using TensorBoard without installing anything locally.

Setup

Start by installing TF 2.0 and loading the TensorBoard notebook extension:

For Jupyter users: If you've installed Jupyter and TensorBoard into the same virtualenv, then you should be good to go. If you're using a more complicated setup, like a global Jupyter installation and kernels for different Conda/virtualenv environments, then you must ensure that the `tensorboard` binary is on your `$PATH` inside the Jupyter notebook context. One way to do this is to modify the `kernel_spec` to prepend the environment's `bin` directory to `$PATH`, [as described here](#).

For Docker users: In case you are running a [Docker](#) image of [Jupyter Notebook server](#) using [TensorFlow's nightly](#), it is necessary to expose not only the notebook's port, but the TensorBoard's port. Thus, run the container with the following command:

```
docker run -it -p 8888:8888 -p 6006:6006 \
tensorflow/tensorflow:nightly-py3-jupyter
```

where the `-p 6006` is the default port of TensorBoard. This will allocate a port for you to run one TensorBoard instance. To have concurrent instances, it is necessary to allocate more ports. Also, pass `--bind_all` to `%tensorboard` to expose the port outside the container.

```
# Load the TensorBoard notebook extension
%load_ext tensorboard
```

Import TensorFlow, datetime, and os:

tensorflow_in_notebooks.ipynb - Colab

https://colab.research.google.com/github/tensorflow/tensorboard/blob/master/docs/tensorboard_in_notebooks.ipynb#scrollTo=TB0wBWfcVqHz&p... 1/7

```
train_model()
```

```
/usr/local/lib/python3.12/dist-packages/keras/src/layers/resizing/Flatten.py:37
super().__init__(**kwargs)
Epoch 1/5
1875/1875 _____ 14s 7ms/step - accuracy: 0.7833 - loss: 0.6109 -
Epoch 2/5
1875/1875 _____ 19s 6ms/step - accuracy: 0.8570 - loss: 0.3917 -
Epoch 3/5
1875/1875 _____ 11s 6ms/step - accuracy: 0.8678 - loss: 0.3526 -
Epoch 4/5
1875/1875 _____ 21s 6ms/step - accuracy: 0.8767 - loss: 0.3333 -
Epoch 5/5
1875/1875 _____ 21s 6ms/step - accuracy: 0.8831 - loss: 0.3124 -
```

Start TensorBoard within the notebook using [magics](#):

```
%tensorboard --logdir logs
```

TensorBoard

TIME SERIES

INACTIVE

Filter runs (regex)

Filter tags (regex)

All

Scalars

Image

Histogram

Run

20250915-052331/train

20250915-052331/val

20250915-052513/train

20250915-052513/val

Pin

20250915-052331/train

Pin cards for a quick view and comparison

bias 2 cards

bias/histogram

20250915-052331/...

bias/histogram

20250915-052513/...

Settings

GENERAL

Horizontal Axis

Step

Enable step selection and data tab (Scalars only)

Enable Range Selection

Link by step 4

Card Width

Enable saving pins (Scalars only)

SCALARS

Smoothing

0.6

Tooltip sorting method

Alphabetical

Ignore outliers in chart scaling

Partition non-monotonic X axis

HISTOGRAMS

You can now view dashboards such as **Time Series**, **Graphs**, **Distributions**, and others. Some dashboards are not available yet in Colab (such as the profile plugin).

The `%tensorboard` magic has exactly the same format as the TensorBoard command line invocation, but with a `%`-sign in front of it.

You can also start TensorBoard before training to monitor it in progress:

%tensorboard --logdir logs

Reusing TensorBoard on port 6006 (pid 1045), started 0:00:01 ago. (Use '!kill 1045' to kill it.)

TensorBoard

TIME SERIES

INACTIVE

Filter runs (regex)

All

Scalars

Image

Histogram

Run

epoch_accuracy

epoch_accuracy

epoch_learning_rate

epoch_loss

evaluation_accuracy_vs_iteration

evaluation_loss_vs_iterations

epoch_accuracy

epoch_learning_rate

epoch_loss

evaluation_accuracy_vs_iteration

evaluation_loss_vs_iterations

Settings

GENERAL

SCALARS

HISTOGRAMS

Horizontal Axis

Step

Enable step selection and data tab (Scalars only)

Enable Range Selection

Link by step 4

Card Width

Smooth

Smoothing

Tooltip sorting method

Ignore outliers in chart scaling

Partition non-monotonic X axis

The same TensorBoard backend is reused by issuing the same command. If a different logs directory was chosen, a new instance of TensorBoard would be opened. Ports are managed automatically.

Start training a new model and watch TensorBoard update automatically every 30 seconds or refresh it with the button on the top right:

```
train_model()

/usr/local/lib/python3.12/dist-packages/keras/src/layers/resizing/flatten.py:37
super().__init__(**kwargs)
Epoch 1/5      1875/1875      13s 6ms/step - accuracy: 0.7795 - loss: 0.6168 -
Epoch 2/5      1875/1875      20s 6ms/step - accuracy: 0.8578 - loss: 0.3876 -
Epoch 3/5      1875/1875      12s 6ms/step - accuracy: 0.8687 - loss: 0.3601 -
Epoch 4/5      1875/1875      12s 6ms/step - accuracy: 0.8783 - loss: 0.3310 -
Epoch 5/5      1875/1875      21s 7ms/step - accuracy: 0.8849 - loss: 0.3126 -
```

You can use the `tensorboard.notebook` APIs for a bit more control:

```
from tensorboard import notebook
notebook.list() # View open TensorBoard instances

Known TensorBoard instances:
- port 6006: logdir logs (started 0:01:26 ago; pid 1045)

# Control TensorBoard display. If no port is provided,
# the most recently launched TensorBoard is used
notebook.display(port=6006, height=1000)
```

Selecting TensorBoard with logdir logs (started 0:01:26 ago; port 6006, pid 10...

TensorBoard

TIME SERIES

INACTIVE

Filter runs

Filter tags (regex)

All

Scalars

Image

Histogram

Run

20250915-052331/t

20250915-052331/v

20250915-052513/t

20250915-052513/v

Pinned

Pin cards for a quick view and comparison

bias 2 cards

bias/histogram

20250915-052331/...

Settings

GENERAL

Horizontal Axis

Step

☒ Enable step selection and data tab (Scalars only)

☐ Enable Range Selection

☐ Link by step 4